

Embodied Neuromorphic Benchmarks

Terrence C. Stewart^{1,*}, Travis DeWolf¹, Ashley Kleinhans² and Chris Eliasmith¹

¹Centre for Theoretical Neuroscience, University of Waterloo, Waterloo, ON, Canada

²Mobile Intelligent Autonomous Systems group, Council for Scientific and Industrial Research, Pretoria, South Africa

Correspondence*:

Terrence C. Stewart

Centre for Theoretical Neuroscience, University of Waterloo, Waterloo, ON, Canada,
tcstewar@uwaterloo.ca

2 ABSTRACT

Evaluating the effectiveness and performance of neuromorphic hardware is difficult. It is even more difficult when the task of interest is an embodied task; that is, a task where the output from the neuromorphic hardware affects its future input through some environment. However, embodied situations are one of the primary potential uses of neuromorphic hardware. To address this, we present a methodology for embodied benchmarking that makes use of a hybrid of real physical embodiment and a type of “minimal” simulation. Minimal simulation has been shown to lead to robust real-world performance, while still maintaining the practical advantages of simulation, such as making it easy for the same benchmark to be used by many researchers. These benchmarks are flexible, in that they allow researchers to explicitly modify the benchmark to identify particular task domains where particular hardware excels. To demonstrate the method, we present a novel benchmark where the task is to perform motor control on an arbitrary system with unknown external forces.

Keywords: neuromorphic hardware, benchmarking, minimal simulation, adaptive control, neural networks

1 INTRODUCTION

Neuromorphic hardware holds great promise for a wide variety of applications. The combination of massively parallel computation and low power consumption means that there is the potential to have complex algorithms running in embedded processing situations, without being a significant drain on available energy. A crucial challenge is to identify what sort of always-on or interactive functionality can best exploit these devices.

To evaluate applications of neuromorphic hardware, we need benchmark tasks. These tasks must allow us to compare across different instances of neuromorphic hardware (and potentially across different algorithms implemented in that hardware). Good benchmarks will allow us to quantitatively compare systems, letting researchers both measure the progress in the field, and also directly compare competing approaches.

In this paper, we focus on the development of *embodied* benchmarks. These are dynamic tasks where the output of the neuromorphic hardware *influences its own future input* through some environment. This is in

contrast to standard categorization or pattern identification tasks, where the input is some fixed sequence and the hardware must produce the correct output for each input (or input pattern).

We believe embodied benchmarks should be of particular interest to neuromorphic research, given that the most compelling applications of neuromorphic hardware are likely to be in this domain of embedded and interactive control of robotic or other physical systems. However, embodiment itself raises a number of issues that complicate the development of such benchmarks. Rather than simply providing a data file of inputs and desired outputs, the benchmark must either specify a full physical system for that embodiment, or it must provide software for a simulation of that system. As we discuss below, both approaches are problematic. Describing a method for overcoming these shortcomings is the primary goal of this paper.

2 EMBODIED BENCHMARKS

An embodied benchmark task is one where the system we are studying has a two-way interaction with some sort of environment. That is, the outputs from the neuromorphic hardware are sent to the environment where they cause an effect, the results of which change the subsequent input. For example, the outputs might control the movement of a robot, which in turn affects the sensory data received by the robot.

2.1 Simulation versus Physical Instantiation

To define such a benchmark, we need to be explicit about the embodiment. If a robot is to be controlled, we need to be explicit about all of the details of that robot. What motors does it have? How are they configured? How strong are they? What sensors are there? Where are they placed? How accurate are they? However, even if these questions are answered, there is a fundamental problem in that *other researchers need access to that exact robot*. If a benchmark is to be widely used, other researchers developing their own neuromorphic hardware should be able to do their own testing on the same benchmark system.

Furthermore, using a physical robot imposes significant practical difficulties when performing extensive benchmark testing. When testing, we often want to run the same task over and over again, both for robustness and to see the effects of varying parameters. With a physical robot, this means manually setting up the task, letting the test run, gathering the resulting data, and then resetting the robot back to the initial state. Consequently, issues like battery life become problematic, and not just because there is a limited amount of time available for testing. As the battery level changes, the performance of the robot itself can also change. Furthermore, for any rigorous testing of the benchmark, we will want to examine situations where the system fails. This means that some of the testing will involve parameter settings that lead to poor behaviour, which might have the undesirable result of causing physical damage.

However, *not* using a real physical embodiment for testing is also problematic. First and foremost, without an actual real-world task, why should we have any confidence that the performance on the benchmark is reflective of the actual usefulness of the neuromorphic hardware? It is widely known that simulations of robots (or other physical systems) are often *much* easier to control and better-behaved than the real thing [TODO: ref?]. The field of robotics is filled with algorithms that work well “in theory,” but fail when run on actual hardware. We do not want a benchmark that falls into this trap, giving high scores to hardware that does not turn out to function well when generalized to real situations.

Furthermore, neuromorphic hardware has another constraint that severely limits the utility of simulation. Typically, when a simulation is too simple to reflect reality, we add details to the simulation itself. Incredibly finely detailed simulations can be created, filling in all of the details needed. However, accurate modelling of physical systems can very quickly become *impractical to run in real time*. This is a fundamental problem,

in that neuromorphic hardware is often tied to real-time interactions, and there is no way to slow down the hardware to match the simulated environment. This means that even if we spent the considerable amount of research effort needed to define a simulated environment for an embodied benchmark, running that simulation fast enough to interact with the desired hardware would require more resources than are available.

2.2 Minimal Simulation

The above considerations seem to indicate that even though using real-world physical hardware for benchmarking is problematic, it is still better than using simplistic simulations which may not generalize to real tasks. However, we believe neuromorphic benchmarking can effectively exploit an approach known as *Minimal Simulation* (Jakobi 1997).

This approach was first suggested in the context of evolutionary robotics. Notably, the problem faced by embodied neuromorphic benchmarking is remarkably similar to that faced earlier by these researchers. In evolutionary robotics, the goal is to use genetic algorithms to *evolve* systems that can control robots to perform various tasks. These tasks can be as simple as navigation and obstacle avoidance, but have also included more difficult tasks such as walking, object identification, and visual tracking [refs??].

However, performing evolution on real physical robots is problematic for the same reasons that benchmarks on physical robots are problematic. The robots must be reset to the same state each time; they often involve behaviour that can physically damage the robots; and they take a very long time to run. For this reason, attempts were made to evolve algorithms using simulated robots. However, the general finding was that algorithms that worked on the simulated robots would not work when run on the real physical robots. If the simulations were improved, adding complex physical detail, then it was possible to generalize to real behaviour; unfortunately, such complex simulations would run slower than real-time [refs].

To address this problem, Jakobi (1997) proposed the creation of “minimal” simulations. These are simulations where there is variability *within the simulation itself*. In other words, we make *poor* simulations, but ensure that the way in which they are poor is itself variable. We are then in a position to ensure that the controllers work across that whole range of variability. “Instead of trying to eliminate the differences between simulation and reality, they are acknowledged, and mechanisms are put in place to prevent evolving controllers from relying on them.” [ref: jakobi thesis].

With this approach, it became possible to build minimal simulations that would run faster than real-time and yet also be complex enough that if a system could successfully control the simulation, it was also likely to successfully control a real robot. To achieve this kind of transfer, the simulations were made to be unreliable in almost every respect. For example, for a simulation of a simple motor it would still be the case that if power is applied it would generally try to spin, but the exact amount of torque, the amount of sensory noise, the amount of time needed, the amount of static and dynamic friction, and so on would all be randomly chosen. A successful controller would have to deal with this wide range of variability, and if it could handle that variability then there would be reason to believe it could also handle the real physical system.

It is worth noting that a minimal simulation only has to be a good simulation *for successful behaviour*. That is, “if we are evolving corridor following behaviour, the dynamics of the simulation might differ wildly from those of reality if the controller hits a wall or goes round in circles, but this does not matter, since the controllers we are interested in transferring across the reality gap will neither hit walls nor go round in circles.” [ref: jakobi thesis] If the controller is poor, we do not need the simulation to be at all

accurate in exactly *how* that poor behaviour is manifest. We do not need an exact detailed physics model of the collision between a robot and a wall, or a detailed model of what happens to a robot arm when it starts oscillating wildly due to a poor control signal. All we need is for the simulation to be just good enough to indicate that things have gone wrong, and thus give a low score to that controller. This means that, for example, in a minimal simulation of an eight-legged walking robot, it is not necessary to have a physics simulation that correctly models what happens when two legs collide with each other. Rather, if legs collide with each other, that is an indication that the walking behaviour is very poor. As long as that result is indicated we can greatly simplify the simulation by not including all the details necessary to deal with these sorts of physical interactions.

2.3 Minimal Simulation as a Benchmark

Given the success of this approach for evolutionary robotics, we propose using a minimal simulation as a neuromorphic benchmark. First, we note that one important use of a benchmark is that by knowing how well particular hardware performs on that benchmark, you can reasonably infer how well that hardware will perform in other situations. For example, if an image recognition algorithm performs well on the MNIST hand-written digit recognition benchmark, we can use that knowledge to guess that it may also perform well on a different recognition task. Of course, this inference will fail if that algorithm has been specifically over-fit to exactly that one situation. For that reason, it would be useful to have a benchmark that covers a large range of variations on the task. If the hardware performs well across that variability, then it is more likely to also work in whatever new situation we want to use it in.

To achieve this, we need software simulations of the environment for the task. These simulations must be fast enough to run in real time (so that they can be controlled by real neuromorphic hardware), and they must be extremely variable. Each time the simulation is run, different parameters will be chosen for this variability (so one run might have a large degree of sensor noise while the next run has none at all; one run might have more delay in the motor response and another might have less power available). Being successful at the benchmark means being successful across all this variability.

The result should be a benchmark that can be run by any researcher. The fact that it is a simulation means that source code can be shared, and that no specialized hardware is needed. Furthermore, the variability in the simulation itself can be controlled, and this can help give a rich characterization of the benchmarked hardware. For example, some hardware might only work with small amounts of sensor noise, or other hardware might only work when there is significant delay in the motor response. This flexibility in parameters in the benchmark allows researchers to explicitly characterize that particular situations where their hardware excels.

2.4 Cheap Robotics

The minimal simulation described above forms the core of our benchmark. However, the point of that benchmark is that it should do a reasonable job of generalizing to real-world physical tasks. It is then useful to supplement that simulation benchmark with at least one easy-to-construct physical analog. This physical version would be one particular instance of the type of situation the benchmark is meant to cover. For that reason, it is much more restrictive in terms of what general conclusions can be drawn from how well different hardware performs in that situation. Rather, it gives an explicit and understandable double-check that models that perform well on the simulation benchmark also perform well in a physical environment. Even considering the advantages of benchmarking using minimal simulations, it is still useful to have a real physical point of comparison as well.

For this physical aspect of the benchmark, we recommend cheap, widely-available components. This allows a greater chance for other researchers to have access to the same (or similar) hardware. For the particular example benchmark described in the next section, we use the Lego Mindstorms EV3 kit, a simple robotics platform available at most toy stores.

It is important to note that there is actually a theoretical advantage to using cheap robotics hardware for benchmarking, in addition to the practical advantages. In particular, we *don't want benchmarks that rely on high-speed, high-accuracy devices*. The purpose of the benchmark is not to indicate how well this neuromorphic hardware works to control this one particular robot in this task. Rather, the purpose of a benchmark is to characterise how well this neuromorphic hardware works on this task *in general*. The variability in the minimal simulation means that it should be able to function across a wide variety of physical embodiments, and so if we are to choose one particular physical embodiment to test in the real world, then we should choose one that is not extremely high-precision. For this reason, we believe using cheap Lego robot is actually more useful for benchmarking than an expensive high-precision robot.¹

3 A BENCHMARK: ADAPTIVE MOTOR CONTROL

To demonstrate this approach to creating embodied neuromorphic benchmarks, we now consider a basic control task. Suppose we have a system with a number of joints q and we want to send an output u to those motors such that the joints move to a particular desired position q_d . Our only output is the signal u (one for each motor) and our only inputs are the current position of each motor q and the desired positions q_d .

The simplest controller for such a situation is a P (proportional) controller, where $u = K_p(q_d - q)$. This is often supplemented with a D (derivative) term, which helps to slow the system down as it approaches the desired position, thus avoiding overshooting and oscillation ($u = K_p(q_d - q) + K_d(\dot{q}_d - \dot{q})$), leading to the standard PD controller. Both K_p and K_d are constants that can be tuned to particular situations.

However, this controller has difficulty in the presence of significant external forces. For example, consider a single motor controlling the angle of a single arm. If the arm is held out to the side, gravity acting on the mass of the arm itself will pull the arm downward. Thus to hold the arm still at a particular q_d will require the controller to apply a force to counteract gravity. Since the PD controller always produces an output $u = 0$ when $q = q_d$, it cannot compensate for this, and so the arm will stay stationary at some angle below the desired angle. [TODO: add diagram]

The standard solution to this problem is to add an I (integral) term ($K_i \int (q_d - q)dt$) to the controller, making it a PID controller. The idea here is that as the difference between where it is and where we want it to be accumulates over time, the K_i term gradually increases how much extra force is being applied until it is large enough to counteract the external force of gravity (or whatever other external forces are present). However, this approach has great difficulty when q_d changes, since the external force due to gravity changes depending on the position of the arm q . The controller ends up having to “relearn” the correct amount of extra force needed every time q_d changes.

In some robotics applications, the solution to this problem is to mathematically analyze the geometry and mass of the system and compute exactly how much extra force is needed. In this particular case, the answer is straight-forward, in that the extra torque due to gravity is $\tau = mg\frac{l}{2}\sin(q)$, where m is the mass of the arm, l is the length, and g is $9.8m/s^2$. If the force applied by the motor is linear in u , then we could simply compute this value and add it to our controller's output. However, this assumes a perfectly even

¹ Of course, for more complex benchmark tasks we may need sensory and motor capabilities that are beyond that of a simple Lego robot.

190 distribution of weight in the arm, ignores momentum and other forces, and gets much more complex as
 191 more joints are added. Furthermore, if this initial computation is slightly off, or if details of the system
 192 change, there is no way to adjust this compensation.

193 Fortunately, there is an adaptive solution to this problem, and it is one that fits well with neuromorphic
 194 hardware. Slotine and Li (1987) shows that if you express the influence of these other external forces
 195 as $\tau = Y(q)\omega$ (where $Y(q)$ is a fixed set of functions of q , such as $\sin(q)$, and ω is a vector of scalar
 196 weights, one for each function in Y), then you can learn to compensate for these external forces by using
 197 the learning rule $\Delta\omega = \alpha Y(q)u$, where u is the basic PD control signal.

198 Importantly, as pointed out by Sanner and Slotine (1992) and Lewis (1996), rather than making explicit
 199 assumptions about the exact functions that should be in $Y(q)$, we can use a neural network approach where
 200 each neuron is a different function of q . As long as there is enough heterogeneity (i.e. as long as the neural
 201 activity forms a basis space that is capable of approximating the external forces), then the learning rule will
 202 continue to work. This approach has been extended to biologically plausible neurons and been used in both
 203 the Recurrent Error-driven Adaptive Control Hierarchy model of human motor control (DeWolf 2014) and
 204 quadcopter control (Komer 2015).

205 This then suggests an explicit neuromorphic benchmark. The input to the neuromorphic hardware is q ,
 206 the system state. This input is fed to each neuron such that each neuron produces some output behaviour
 207 that is based on this input. Since q will be multi-dimensional (if there is more than one joint), we may give
 208 each neuron a random weighting of each q value ($J_i = e_i \cdot q$, where J_i is the input to neuron i , and e_i is
 209 a randomly chosen vector²). Given this input, the neurons will produce some output A . We now form a
 210 weighted sum of these outputs Ad , where d is a matrix (number of neurons by number of elements in q)
 211 that is initially all zeros.

212 To use this controller, we add its output to that of the standard PD controller. That is, the standard
 213 controller has $u = K_p(q_d - q) + K_d(\dot{q}_d - \dot{q})$, and so our actual output to the motor is $u + Ad$. We then
 214 apply a learning rule on d such that $\Delta d = \alpha A \times u$.

215 Notice that we can think of this system as a three-layer neural network, where the input and output layers
 216 are linear. The first layer is q , the input state, one value for each joint. The “hidden” layer is A , the activity
 217 of a large number of neurons. The output layer again has one value per joint, and is the extra added signal
 218 to apply to the motors, Ad . Given that this is such a canonical example of the use of neural networks, we
 219 hope that the majority of neuromorphic hardware is flexible enough to implement exactly this model.

220 3.1 Online and offline learning

221 The one major step here that does not exist in a lot of neuromorphic hardware is the ability to update the
 222 weights d . For hardware that does have a built-in learning rule, this rule is at least of a very common form,
 223 where the weight update from a neuron is proportional to the activity of that neuron and an external error
 224 signal. This makes it an instance of the ubiquitous delta rule, and hopefully supported by the hardware.

225 However, if the neuromorphic hardware being benchmarked does not have the ability to update weights
 226 online using a learning rule of this form, then there are two solutions. First, the multiplication by d could be
 227 done on the output from the neuromorphic hardware. There has to be some system to take the neural output
 228 from the hardware and send it to the motor (or to the simulation of the motor). Instead of outputting the

² e could also be chosen so as to regularly span the space of possibilities

result of Ad , the hardware could output A (the activity of all the neurons), and the interface to the motor could be responsible for doing the multiplication by d and updating d according to the learning rule.

Alternatively, we can use offline learning. That is, rather than updating the weights d all the time, we simply record A and u , and then after a period of time stop the controller, compute the sum total of the changes to d , load the new value of d onto the neuromorphic hardware, and then start the controller again.

3.2 Minimal Simulation for Adaptive Control

Now that we have defined the task, we can use the principles of Minimal Simulation to construct a flexible and variable simulated environment for testing adaptive control. The idea is to make a bare-bones simulation of the system being controlled, with built-in variability. If the neuromorphic controller works well across this variability, then it is likely to work well outside of simulation as well.

The basic system variable is a vector of length N holding the joint angles q . Each joint has a velocity v . The force applied by each motor is related to the signal sent to the motor u , but will generally have some maximum value T , so we use $\tanh(u)T$. For friction, we simply scale the velocity by some factor F every time step. This results in the simplistic simulation as follows:

$$\Delta v = -vF + \tanh(u)T \quad (1)$$

$$\Delta q = v \quad (2)$$

On top of this, we need to add an external perturbing force. In a real system, this could be the effects of gravity given the current configuration of the motors, or of other unexpected influences. Rather than choosing one particular fixed external force for our benchmark, we *randomly generate* this force each time the benchmark is run.

We start with a small set of smooth functions f which are often found in dynamics equations (e.g. x , x^2 , $\sin(x)$). We then generate an external force of $K_f(\zeta \cdot f(\beta \cdot q + \gamma) + \eta)$ where ζ , β , γ , and η are all random vectors and K_f is a scaling factor to control how strong this external force is. The result is added to Equation 1. Note that this means that if q is 4-dimensional (i.e. if there are four joints being controlled) and if there are three smooth functions in f (as there are here), then β , γ , and η are all vectors of length 4 and ζ is a 4x12 matrix. All of these values are randomly chosen from the normal distribution $N(0, 1)$.

Finally, we add random noise, delay, and filtering to both the input and the output of the system. For noise, we add $N(0, \sigma_u)$ to the control signal u and $N(0, \sigma_q)$ to the q value reported back to the controller. We also use a low-pass filter to smooth both values (with time constants τ_u and τ_q) after this noise is added. Finally, both q and u are delayed by an amount of time t_q and t_u .

The result is not meant to be a simulation of a particular physical embodiment. Rather, the variability in this simulation is meant to be extremely fast to simulate, and to make it hard to “cheat” to control it. In other words, if a controller manages to be able to control the various randomly created minimal simulations of embodiment that are generated with this approach, then we have reason to believe that it will also be successful at controlling real embodiments. With this simulation and modern computers, we can easily simulate in real-time systems with dozens of joints and highly complex interactions between them, and measure how well an adaptive controller deals with these situations.

264 3.3 Cheap Robotics for Adaptive Control

265 Of course, we also want an indication that the minimal simulation defined in the previous section
 266 is reasonably representative of the sorts of real-world situations in which we might want to use these
 267 controllers. Importantly, this physical instantiation does not have to exactly match one particular parameter
 268 setting of the minimal simulation. Rather, we want a physical system that shares basic functional similarities
 269 to the minimal simulation defined previously.

270 For example, we want the inputs to the system to act like u , in that a positive number will increase some
 271 velocity v which will in turn increase some sensor value q . We want there to be some sort of external force
 272 applied that affects q , and we want that external force itself to be a function of q . We want there to be
 273 communication delays and noise in the sensory and motor system, and we want all of these effects to be
 274 somewhere within the extreme ranges covered by the minimal simulation. While this cannot prove that
 275 hardware that is successful in simulation will always be successful in any similar real-world task, it at least
 276 gives an existence proof that there is at least one real-world task where it also performs reasonably.

277 For this demonstration, we define an easy-to-build example of a system that can be usefully controlled
 278 by this adaptive method. In particular, we use the Lego Mindstorms EV3 robot kit, organized as shown
 279 in [TODO: add figure]. It consists of a single motor, mounted in such a way that there is a significant
 280 amount of weight on the arm itself (the weight of the motor itself). Multiple motors can be added, and
 281 other configurations can be considered and should also be suitable for testing control, but here we just
 282 consider the basic case.

283 To interface to the physical hardware, we installed the `ev3dev` operating system (<http://ev3dev.org>),
 284 a Debian-based Linux system specifically developed for the EV3. We then installed and ran the `ev3_link`
 285 program from `ev3dev-c` (<https://github.com/in4l1o/ev3dev-c>). This allows the EV3 to
 286 listen for UDP commands that tell it to set motor values and read sensor values. Communication with
 287 a PC was over a USB link (although the system also supports WiFi communication). With constant
 288 communication, the system is able to adjust the power sent to the motors u and give position feedback q
 289 from those motors at a rate of around 200Hz.

290 Figure 1 shows the effects of adaptive control on this physical system. Without adaptation (i.e. with a
 291 simple PD controller), there system state q (the joint angle) overshoots the desired q_d . This overshoot is
 292 largest when q is large. This is because the external force applied to the joint due to gravity is proportional
 293 to $\sin(q)$. The q value also overshoots and comes back part-way, due to physical momentum.

294 However, with adaptation (the right-hand side of Figure 1), the system learns to counteract this extra
 295 force due to gravity. After the first 5 seconds, the system is able to bring q much closer to the desired q_d .

296 Now that we have this physical example of the task our minimal simulation benchmark is meant to cover,
 297 we can use it to calibrate the parameters of the simulation. For example, communication with the EV3
 298 happens around 200Hz, meaning that there must be a delay on the order of 0.005 seconds. Given this, we
 299 set the delays in the simulation to be uniformly chosen between 0 and 0.01. Importantly, we do not need to
 300 exactly measure the delay in the EV3 robot — we just make sure that the minimal simulation is worse.

301 For sensor noise, we note that the EV3 rotation encoders for the motors (the devices that measure q) have
 302 a resolution of 0.0175 (1 degree). This is a very different sort of noise than the gaussian noise used in the
 303 simulation, so we set the simulation noise to be much larger (uniformly distributed between 0 and 0.1).
 304 Similarly, the motor resolution is 0.01, as it accepts integer values up to 100, so we set the motor noise to
 305 be uniform between 0 and 0.1.

306 Finally, we can use the physical system to calibrate the relationship between T (the maximum torque
 307 applied by the motor) and K_f (the scaling factor of the external force). After all, we do not want external
 308 forces that are so strong that the system does not have enough strength to counteract them. On the physical
 309 robot, in the worst-case scenario ($q = \pi/2$ or $-\pi/2$), the motors must be driven at around 0.3 times their
 310 maximum strength to balance the force of gravity. (This can, of course, be adjusted by changing the weight
 311 and its position on the end of the arm). If we arbitrarily fix K_f to 1 and randomly generate external forces
 312 given the process described above, then 95% of the time we get values between -3.75 and +3.75. Since we
 313 want the motors to be strong enough to compensate for forces in that range, we set T to 10.

4 BENCHMARK ANALYSIS

314 - what do we measure? performance: rmse and delay. Also, power consumption? cost?
 315 - for given hardware, how many neurons can be done in realtime? that's how many should be used for
 316 comparison
 317 - show how these things change as we change parameters in the simulation
 318 - delays
 319 - number of motors N
 320 - magnitude of nonlinearity K_f
 321 CPU, GPU, and SpiNNaker?

5 OTHER BENCHMARKS

322 (sketch out how this approach could be applied to other task)
 323 Adaptive Jacobian
 324 Driving down a corridor, avoiding obstacles
 325 Classical conditioning
 326 Operant conditioning

DISCLOSURE/CONFLICT-OF-INTEREST STATEMENT

327 The authors declare that the research was conducted in the absence of any commercial or financial
 328 relationships that could be construed as a potential conflict of interest.

ACKNOWLEDGMENTS

329 Text
 330 Text Text Text Text Text. Text Text Text Text Text Text Text Text Text Text Text Text Text Text Text
 331 Text Text Text Text Text Text Text Text Text Text.
 332 *Funding:* Text Text Text Text Text Text Text Text.

REFERENCES

- 333 DeWolf, T. (2014). *A neural model of the motor control system*. Phd thesis, University of Waterloo
- 334 Jakobi, N. (1997). Evolutionary robotics and the radical envelope-of-noise hypothesis. *Adaptive Behavior*
- 335 6, 325–368
- 336 Komer, B. (2015). *Biologically Inspired Adaptive Control of Quadcopter Flight*. Masters thesis, University
- 337 of Waterloo
- 338 Lewis, F. (1996). Neural network control of robot manipulators. *IEEE Expert: Intelligent Systems and*
- 339 *Their Applications* 11, 64–75
- 340 Sanner, R. and Slotine, J.-J. (1992). Gaussian networks for direct adaptive control. *IEEE Transactions on*
- 341 *Neural Networks* 3
- 342 Slotine, J.-J. E. and Li, W. (1987). On the adaptive control of robot manipulators. *Int. J. Robotics Research*
- 343 6, 49–59. doi:10.1177/027836498700600303

FIGURES

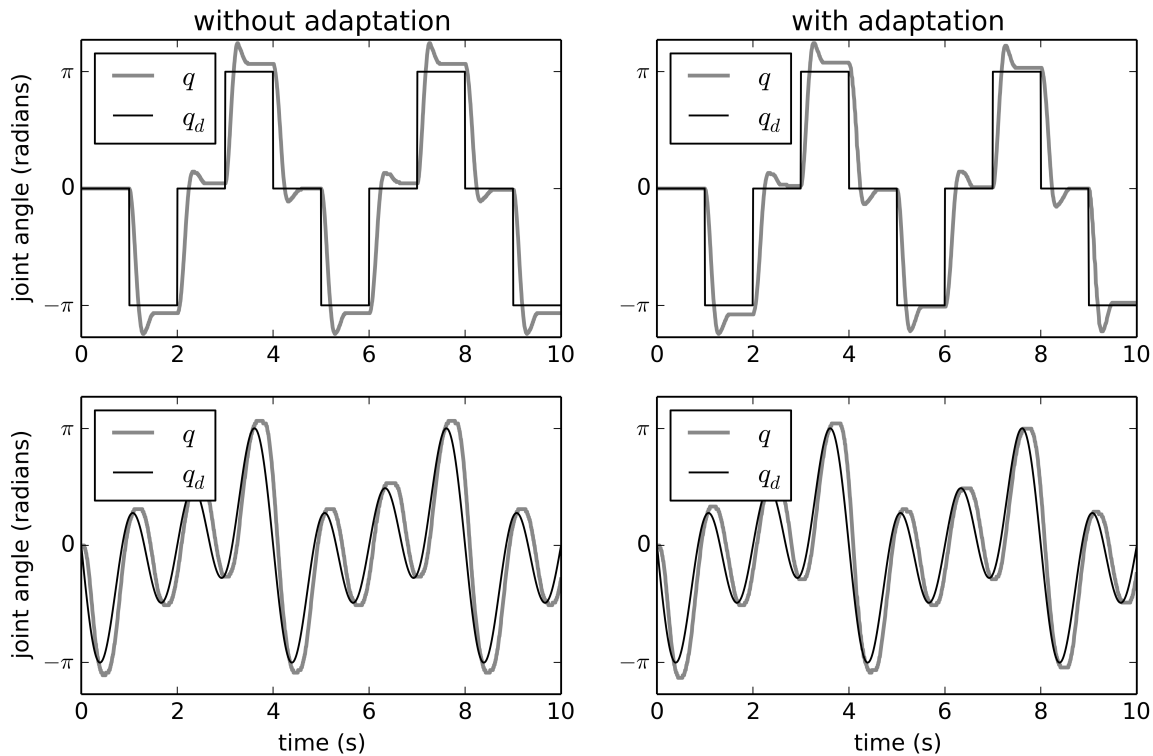


Figure 1. Adaptive control of the EV3 lego robot used for calibrating the minimal simulation. The effects of adaptation over two different desired trajectories are shown. Without adaptation, the joints q do not reach the desired q_d when q_d is large (which is when the external force is largest). With adaptation, q is closer to q_d after about 5 seconds, showing that the system has quickly learned to compensate.